

FinalFunnel - Complete Deployment Guide

System Overview

This system is designed to handle **1TB+ data** with:

- **Sub-second query response** using materialized views and advanced indexing
 - **Bulk upload processing** at 1000-5000 rows/second
 - **PostgreSQL 15** with performance optimization
 - **Redis caching** for ultra-fast filter options
 - **Automated backups** and monitoring
-

Hardware Requirements

Minimum (Testing)

- **CPU:** 4 cores
- **RAM:** 8 GB
- **Storage:** 100 GB SSD
- **Network:** 100 Mbps

Recommended (Production)

- **CPU:** 8+ cores
- **RAM:** 32 GB
- **Storage:** 500 GB NVMe SSD
- **Network:** 1 Gbps
- **Backup Storage:** 1 TB

Large Scale (1TB+ data)

- **CPU:** 16+ cores
- **RAM:** 64 GB+
- **Storage:** 2 TB NVMe SSD (RAID 10)
- **Network:** 10 Gbps
- **Backup Storage:** 5 TB

Quick Start Deployment

Step 1: Prepare Ubuntu Server (20.04 LTS or 22.04 LTS)

```
bash

# Update system
sudo apt update && sudo apt upgrade -y

# Download setup script
wget https://raw.githubusercontent.com/YOUR_REPO/finalfunnel/main/setup_system.sh

# Make executable
chmod +x setup_system.sh

# Run setup (this will take 10-15 minutes)
sudo bash setup_system.sh
```

Step 2: Configure Environment

```
bash

cd /home/ubuntu/finalfunnel

# Copy template and edit
cp .env.template .env
nano .env
```

Edit these critical values:

```
env

# Database - CHANGE PASSWORD!
DB_PASSWORD=YOUR_SECURE_DB_PASSWORD_HERE

# Application - CHANGE PASSWORD!
APP_USERNAME=admin
APP_PASSWORD=YOUR_SECURE_ADMIN_PASSWORD_HERE

# Redis (optional - keep default if running locally)
REDIS_HOST=localhost
```

Step 3: Initialize Database

```
bash
```

```
# Activate virtual environment
```

```
source venv/bin/activate
```

```
# Run database initialization
```

```
python init_database.py
```

Expected output:

```
Creating optimized database schema...
```

```
✓ Database initialized successfully!
```

```
📊 Tables: 0 rows in fact_contacts, 0 in fact_companies
```

```
🔍 Indexes: 45 total indexes created
```

Step 4: Start Application

```
bash
```

```
# Using supervisor (recommended)
```

```
sudo supervisorctl start finalfunnel
```

```
# OR manual start for testing
```

```
streamlit run 0_Home.py --server.port 8501
```

Step 5: Access Application

Open your browser:

- **Local:** <http://localhost:8501>
- **Remote:** <http://your-server-ip:8501>

Default Login:

- Username: (from .env APP_USERNAME)
- Password: (from .env APP_PASSWORD)

🔒 Security Setup

1. Change Default Passwords

PostgreSQL:

```
bash
```

```
sudo -u postgres psql  
ALTER USER finalfunnel_user WITH PASSWORD 'your_new_strong_password';  
\q
```

Update `.env` with new password.

2. Setup SSL with Let's Encrypt

```
bash
```

```
# Edit nginx config first  
sudo nano /etc/nginx/sites-available/finalfunnel  
  
# Change: server_name your-domain.com;  
  
# Get SSL certificate  
sudo certbot --nginx -d your-domain.com  
  
# Auto-renewal is setup automatically
```

3. Firewall Configuration

```
bash
```

```
# Allow SSH, HTTP, HTTPS  
sudo ufw allow 22/tcp  
sudo ufw allow 80/tcp  
sudo ufw allow 443/tcp  
  
# Block direct access to Streamlit port  
sudo ufw deny 8501/tcp  
  
# Enable firewall  
sudo ufw enable
```

4. Database Security

```
bash
```

```
# Edit PostgreSQL config
sudo nano /etc/postgresql/15/main/pg_hba.conf

# Change this line:
# local all all peer
# TO:
# local all all md5

# Restart PostgreSQL
sudo systemctl restart postgresql
```

How to Upload Data

Method 1: Web Interface (Up to 500MB)

1. Navigate to **Upload Data** page
2. Click **Download Template**
3. Fill template with your data
4. Click **Upload Your Data**
5. Select file and click **Validate**
6. If validation passes, click **Process Upload**

Method 2: Direct Database Import (500MB+)

For very large files:

```
bash

# 1. Copy file to server
scp your_large_file.csv ubuntu@your-server:/tmp/

# 2. SSH to server
ssh ubuntu@your-server

# 3. Run import script
cd /home/ubuntu/finalfunnel
source venv/bin/activate
python scripts/bulk_import.py /tmp/your_large_file.csv
```

Data Quality Requirements

Required Columns (exact order):

1. comp_name
2. comp_domain
3. annrev
4. comp_industry
5. comp_linkedin
6. firstname
7. lastname
8. jobtitle
9. manlevel
10. empemail (REQUIRED - must be unique)
11. emplinkedin
12. country_code
13. comp_phone
14. comp_street
15. comp_city
16. comp_state
17. comp_country
18. comp_zipcode
19. qa_disposition
20. empsize

Data Formats:

- **empemail:** Must be valid email (required)
 - **annrev:** "10M", "5B", or number
 - **empsize:** "50-200", "1000+" or number
 - **URLs:** No http://, https://, or www. needed
 - **Phone:** Any format accepted
-

🔍 Fast Filtering & Performance

Understanding the Architecture

The system uses a **3-tier data structure**:

1. Fact Tables (fact_contacts, fact_companies)

- Normalized storage
- Optimized for writes
- Foreign key relationships

2. Materialized View (mv_contacts_data)

- Pre-joined denormalized data
- Optimized for reads
- Refreshed after uploads

3. Indexes (45+ specialized indexes)

- B-tree for exact matches
- GIN for text search
- Composite for filter combinations

Query Performance

Expected response times (after warm-up):

Records	Filters	Response Time
1M	0-2	<100ms
10M	0-2	<200ms
100M	0-2	<500ms
1B	0-2	<2s
1M	5+	<50ms
100M	5+	<300ms

Performance Optimization Tips

1. **Use multiple filters** - More filters = faster queries (counter-intuitive but true due to indexes)
2. **Refresh materialized view after bulk uploads:**

sql

```
REFRESH MATERIALIZED VIEW CONCURRENTLY mv_contacts_data;
```

3. Update statistics weekly:

```
sql
```

```
ANALYZE mv_contacts_data;
```

4. Monitor slow queries:

```
sql
```

```
SELECT * FROM pg_stat_statements  
WHERE mean_exec_time > 1000  
ORDER BY mean_exec_time DESC;
```

Maintenance

Daily Tasks (Automated)

- Cache refresh (Redis)
- Backup validation
- Log rotation

Weekly Tasks

```
bash
```

```
# Update statistics  
sudo -u postgres psql -d finalfunnel_db -c "VACUUM ANALYZE;"  
  
# Check index usage  
sudo -u postgres psql -d finalfunnel_db -c "SELECT * FROM v_index_usage WHERE idx_scan < 100;"
```

Monthly Tasks

```
bash
```

```
# Full system backup  
sudo -u postgres pg_dump finalfunnel_db > backup_$(date +%Y%m%d).sql  
  
# Compress  
gzip backup_$(date +%Y%m%d).sql  
  
# Upload to S3 or backup storage
```

Database Size Monitoring

sql

```
-- Check table sizes
SELECT * FROM v_table_sizes;

-- Check growth
SELECT
    schemaname,
    tablename,
    pg_size.pretty(pg_total_relation_size(schemaname||'.'||tablename)) AS size,
    pg_size.pretty(pg_total_relation_size(schemaname||'.'||tablename) -
                  pg_relation_size(schemaname||'.'||tablename)) AS index_size
FROM pg_tables
WHERE schemaname = 'public'
ORDER BY pg_total_relation_size(schemaname||'.'||tablename) DESC
LIMIT 10;
```

Troubleshooting

Application Won't Start

```
bash

# Check logs
tail -f /var/log/finalfunnel.err.log
tail -f /var/log/finalfunnel.out.log

# Check supervisor status
sudo supervisorctl status finalfunnel

# Restart
sudo supervisorctl restart finalfunnel
```

Slow Queries

```
bash
```

```
# Check active queries
sudo -u postgres psql -d finalfunnel_db -c "
SELECT pid, now() - query_start as duration, query
FROM pg_stat_activity
WHERE state = 'active' AND now() - query_start > interval '5 seconds';
"

# Kill slow query
sudo -u postgres psql -d finalfunnel_db -c "SELECT pg_terminate_backend(PID_HERE);"
```

Database Connection Issues

```
bash
```

```
# Check PostgreSQL status
sudo systemctl status postgresql

# Check connections
sudo -u postgres psql -c "SELECT count(*) FROM pg_stat_activity;"

# Increase max connections (if needed)
sudo nano /etc/postgresql/15/main/postgresql.conf
# Change: max_connections = 200
sudo systemctl restart postgresql
```

Out of Disk Space

```
bash
```

```
# Check disk usage
df -h

# Find large files
du -sh /var/lib/postgresql/15/main/*

# Clean old backups
find /home/ubuntu/backups -name "*.sql.gz" -mtime +30 -delete

# Vacuum database
sudo -u postgres psql -d finalfunnel_db -c "VACUUM FULL;"
```



Monitoring & Alerts

Setup Monitoring Dashboard

```
bash

# Install pgAdmin 4 (optional)
curl https://www.pgadmin.org/static/packages_pgadmin_org.pub | sudo apt-key add
sudo sh -c 'echo "deb https://ftp.postgresql.org/pub/pgadmin/pgadmin4/apt/$(lsb_release -cs) pgadmin4 main" > /etc/apt/sources.list.d/pgadmin4.list'
sudo apt update
sudo apt install pgadmin4
```

Key Metrics to Monitor

1. Database Size Growth
2. Query Response Time
3. Active Connections
4. CPU & Memory Usage
5. Disk I/O
6. Upload Success Rate

Backup & Recovery

Automated Backup Script

Create `/home/ubuntu/finalfunnel/scripts/backup.sh`:

```
bash

#!/bin/bash
BACKUP_DIR="/home/ubuntu/backups"
DATE=$(date +%Y%m%d_%H%M%S)
DAYS_TO_KEEP=30

# Create backup
sudo -u postgres pg_dump finalfunnel_db | gzip > $BACKUP_DIR/backup_$DATE.sql.gz

# Upload to S3 (optional)
# aws s3 cp $BACKUP_DIR/backup_$DATE.sql.gz s3://your-bucket/backups/

# Delete old backups
find $BACKUP_DIR -name "backup_*sql.gz" -mtime +$DAYS_TO_KEEP -delete

echo "Backup completed: backup_$DATE.sql.gz"
```

Setup cron:

```
bash

crontab -e

# Add: Daily backup at 2 AM
0 2 * * * /home/ubuntu/finalfunnel/scripts/backup.sh >> /var/log/backup.log 2>&1
```

Recovery Process

```
bash

# Stop application
sudo supervisorctl stop finalfunnel

# Drop and recreate database
sudo -u postgres psql -c "DROP DATABASE finalfunnel_db;" 
sudo -u postgres psql -c "CREATE DATABASE finalfunnel_db;" 

# Restore from backup
gunzip -c backup_20240115_020000.sql.gz | sudo -u postgres psql finalfunnel_db

# Restart application
sudo supervisorctl start finalfunnel
```

📞 Support & Resources

Log Locations

- Application: (/var/log/finalfunnel.out.log)
- Errors: (/var/log/finalfunnel.err.log)
- PostgreSQL: (/var/log/postgresql/postgresql-15-main.log)
- Nginx: (/var/log/nginx/access.log) & (error.log)

Useful Commands

```
bash

# Application status
sudo supervisorctl status finalfunnel

# Database size
sudo -u postgres psql -d finalfunnel_db -c "SELECT pg_size.pretty(pg_database_size('finalfunnel_db'));" 

# Active users
sudo -u postgres psql -d finalfunnel_db -c "SELECT count(*) FROM pg_stat_activity WHERE datname='finalfunnel_db';"

# Clear cache
redis-cli FLUSHALL

# Restart everything
sudo supervisorctl restart finalfunnel
sudo systemctl restart postgresql
sudo systemctl restart redis-server
sudo systemctl restart nginx
```

⌚ Best Practices

- 1. Regular Backups** - Automate daily backups
- 2. Monitor Performance** - Setup alerts for slow queries
- 3. Update Statistics** - Weekly ANALYZE
- 4. Security Updates** - Monthly OS updates
- 5. Test Restores** - Quarterly backup restoration tests
- 6. Capacity Planning** - Monitor growth trends
- 7. Documentation** - Keep deployment notes
- 8. Access Control** - Limit database access
- 9. SSL/TLS** - Always use HTTPS in production
- 10. Change Passwords** - Never use defaults

✓ Post-Deployment Checklist

- Ubuntu server updated
 - PostgreSQL installed and configured
 - Database initialized with schema
 - Application running via supervisor
 - Nginx configured with SSL
 - Firewall rules applied
 - Passwords changed from defaults
 - Automated backups setup
 - Monitoring configured
 - Test upload completed
 - Filter performance verified
 - Documentation reviewed
 - Admin access confirmed
 - Support contacts saved
-

System Version: 1.0.0

Last Updated: November 2024

Designed for: 1TB+ data with sub-second response times